

# Lab 3: Descriptive Statistics

[Code](#)

AUTHOR

Maxim Dokukin

PUBLISHED

February 13, 2024

**Follow the instructions below and use R Markdown to create a pdf document with your code and answers to the following questions on Gradescope.** You may find a template file by clicking "Code" in the top right corner of this page.

Your final submission should clearly include all code needed to generate your answers and should be formatted according to the guidelines outlined in class. In particular, make sure:

1. Code and output are clearly organized by question.
2. Unnecessary messages, warning, and output are removed.

You may collaborate with your classmates and consult external resources, but you should write and submit your own answer. **Any classmates with whom you collaborate should be credited at the top of your submission. Similarly, if you consult any external references, you should cite them clearly and explicitly.**

## A. Weather Forecast Data

1. For this lab, we'll be using data on weather forecasts gathered by student at Saint Louis University. You can read about the dataset [here](#). Download the weather forecasts data using the following code:

```
weather_forecasts <- readr::read_csv('https://raw.githubusercontent.com/slsu-weather-forecasts/summary(weather_forecasts)
```

	date	city	state
high_or_low			
Min.	:2021-01-30	Length:651968	Length:651968
Length:651968			
1st Qu.:	2021-05-31	Class :character	Class :character
Class :character			
Median :	2021-09-30	Mode :character	Mode :character
Mode :character			
Mean	:2021-09-30		

3rd Qu.:2022-01-30

Max. :2022-06-01

forecast_hours_before	observed_temp	forecast_temp	observed_precip
Min. :12	Min. : -47.00	Min. : -41.00	Min. : 0.00
1st Qu.:21	1st Qu.: 42.00	1st Qu.: 42.00	1st Qu.: 0.00
Median :30	Median : 59.00	Median : 59.00	Median : 0.00
Mean :30	Mean : 57.56	Mean : 57.36	Mean : 0.10
3rd Qu.:39	3rd Qu.: 74.00	3rd Qu.: 74.00	3rd Qu.: 0.02
Max. :48	Max. :122.00	Max. :118.00	Max. :12.40
	NA's :47744	NA's :37313	NA's :50416

forecast\_outlook possible\_error  
 Length:651968 Length:651968  
 Class :character Class :character  
 Mode :character Mode :character

2. How many rows are in this dataset? How many columns?

```
dim(weather_forecasts)
```

```
[1] 651968    10
```

651968 rows, 10 cols.

3. How many cities are represented in this dataset?

```
length(unique(weather_forecasts$city))
```

```
[1] 160
```

160 cities = size of array with unique cities.

4. Create a new data frame containing only the forecasts for San Jose.  
You may have to explore the values for the `city` variable.

```
san_jose_data = weather_forecasts[weather_forecasts$city == 'SAN JOSE']
```

15 min wasted on realizing its `san_jose`, not `san jose`.

5. Compute the mean absolute error between `observed_temp` and `forecast_temp` for San Jose.

```
san_jose_data$absolute_error <- abs(san_jose_data$observed_temp -
                                     san_jose_data$forecast_temp)
mean_error <- mean(san_jose_data$absolute_error, na.rm = TRUE)
print(mean_error)
```

```
[1] 2.169762
```

2.169762

6. Compute the mean absolute error between `observed_temp` and `forecast_temp` for San Jose using only forecasts made 48 hours in advance.

```
san_jose_48adv_data = san_jose_data[san_jose_data$forecast_hours == 48]
san_jose_48adv_data$absolute_error <- abs(san_jose_48adv_data$observed_temp -
                                           san_jose_48adv_data$forecast_temp)
mean_error_48 <- mean(san_jose_48adv_data$absolute_error, na.rm = TRUE)
print(mean_error_48)
```

```
[1] 2.262544
```

2.262544

7. Compute the mean absolute error between `observed_temp` and `forecast_temp` for San Jose using only forecasts made 12 hours in advance.

```
san_jose_12adv_data = san_jose_data[san_jose_data$forecast_hours == 12]
san_jose_12adv_data$absolute_error <- abs(san_jose_12adv_data$observed_temp -
                                           san_jose_12adv_data$forecast_temp)
mean_error_12 <- mean(san_jose_12adv_data$absolute_error, na.rm = TRUE)
```

```
print(mean_error_12)
```

```
[1] 2.0553
```

```
2.0553
```

8. Compare your answers to 6 and 7. What do you notice? How does this compare to your expectation?

24h advance forecasts have smaller error. This makes sense, as it is easier to make more accurate predictions in the short term.

9. Pick two cities in this dataset. Investigate whether the forecast accuracy is better for one city than for the other, using an appropriate statistic. Discuss your findings.

```
nyc_data = weather_forecasts[weather_forecasts$city == 'NEW_YORK']
nyc_data$absolute_error <- abs(nyc_data$observed_temp - nyc_data$forecasted_temp)
nyc_mean_error <- mean(nyc_data$absolute_error, na.rm = TRUE)
print(nyc_mean_error)
```

```
[1] 2.182927
```

```
gf_data = weather_forecasts[weather_forecasts$city == 'GREAT_FALLS']
gf_data$absolute_error <- abs(gf_data$observed_temp - gf_data$forecasted_temp)
gf_mean_error <- mean(gf_data$absolute_error, na.rm = TRUE)
print(gf_mean_error)
```

```
[1] 3.05578
```

Great Falls accuracy is significantly lower than in NYC. Maybe there are more variables that effect weather there? Yet, the best advice to avoid rain is the same, 'Take the umbrella'

## B. Find your own data

For this component, pick a [Tidy Tuesday dataset](#) and complete the following activity.

10. Provide a brief description of your dataset. Identify at least two questions you could try to answer using this dataset.

```
library(tidyverse)

big_tech_stock_prices <- readr::read_csv('https://raw.githubusercontent.com/

summary(big_tech_stock_prices)
```

```
  stock_symbol      date      open
high
  Length:45088      Min.   :2010-01-04      Min.   : 1.076      Min.
: 1.109
  Class :character  1st Qu.:2013-05-30      1st Qu.: 25.670      1st
Qu.: 25.930
  Mode  :character  Median :2016-08-09      Median : 47.930
Median : 48.460
                                Mean   :2016-08-03      Mean   : 89.267      Mean
: 90.370
                                3rd Qu.:2019-10-21      3rd Qu.:128.662      3rd
Qu.:129.849
                                Max.    :2023-01-24      Max.    :696.280      Max.
:700.990
      low      close      adj_close
volume
  Min.   : 0.9987      Min.   : 1.053      Min.   : 1.053      Min.
:5.892e+05
  1st Qu.: 25.3600      1st Qu.: 25.660      1st Qu.: 22.076      1st
Qu.:9.629e+06
  Median : 47.4650      Median : 47.970      Median : 45.377      Median
:2.646e+07
  Mean   : 88.1119      Mean   : 89.271      Mean   : 85.210      Mean
:5.298e+07
  3rd Qu.:127.2539      3rd Qu.:128.641      3rd Qu.:113.672      3rd
Qu.:5.840e+07
  Max.    :686.0900      Max.    :691.690      Max.    :691.690      Max.
:1.881e+09
```

Stock prices from 2014 to 2022 for big tech.

- What was the most growing stock from Jan 10 2014 to Jan 10 2022 (%)?
- What was the least growing stock from Jan 10 2012 to Jan 10 2022 (%)?
- What is the average (median) growth from Jan 10 2014 to Jan 10

2022 (%)?

PS I was trying Jan 1st, but realized stocks don't trade this day because of the holiday. Jan 10 2015 had no data either. Meta has not been on data set in 2010... had to adjust my date range to make sure I have all the numbers.

11. Open your dataset in R and compute one or more descriptive statistics that shed light on your questions. Discuss your findings.

```
companies <- unique(big_tech_stock_prices$stock_symbol)
prices_2014 <- big_tech_stock_prices[big_tech_stock_prices$date == '2014-01-01', ]
prices_2022 <- big_tech_stock_prices[big_tech_stock_prices$date == '2022-01-01', ]

prices_change <- data.frame(
  companies,
  prices_2014$close,
  prices_2022$close
)

prices_change$percent_change = (prices_change$prices_2022.close -
                                prices_change$prices_2014.close) /
                                prices_change$prices_2014.close * 100
colnames(prices_change) <- c('company', 'price_2014', 'price_2022', 'percent_change')
```

Best performer 2014-2022

```
prices_change[which.max(prices_change$percent_change), ]
```

	company	price_2014	price_2022	percent_change
12	NVDA	3.9325	274	6867.578

Worst performer 2014-2022

```
prices_change[which.min(prices_change$percent_change), ]
```

	company	price_2014	price_2022	percent_change
7	IBM	179.0249	135.03	-24.57472

Median growth rate 2014 - 2022

```
median(prices_change$percent_change)
```

```
[1] 589.2025
```

12. Are there any limitations of your analysis? Could additional data or more complicated methods improve your analysis? Discuss.

I wanted to do Jan 1 2010 to Jan 1 2020. But, for Jan 1 there is no data. META also had no data for the early years. I had to adjust my date range to fit these limitations. Using the last value in the previous year before Jan 1, could have made possible data analysis in the range Jan 1, YYYY - Jan 1, YYYY